


# AI & Generative AI in Industry (2025)

Explore Key Trends, Practical Use-Cases, and Career Paths — An Interactive Session for All (No Coding Required)

Karthikeyan NG

Madurai 2025



# Welcome to the Future

Today, we'll dive into the real-world impact of AI on industries, uncover practical projects you can immediately apply, and guide you through navigating career opportunities in this dynamic field. Basic coding experience is needed and your curiosity and a readiness to explore.



# Today's Plan

01

---

## Goals & Interests

Share what brought you here and what you hope to achieve (5 minutes)

03

---

## AI Projects You Can Build Today

Discover practical project ideas for campus life, coding, and creative fields (15 minutes)

05

---

## Group Activity: Design Real-World AI Solutions

Collaboratively create AI solutions for everyday problems (10 minutes)

02

---

## AI Landscape: Breakthroughs & Accessibility

Explore key moments that made AI mainstream and usable today (10 minutes)

04

---

## Interactive: Map Your AI Career

Explore personalized roles, skills, and pathways in the AI industry (15 minutes)

06

---

## Share Your Ideas & Next Steps

Present your solutions and discover resources for continued learning (5 minutes)

# What Changed (Why Now)?

## Broadly Capable AI Models

AI models are now trained on vast amounts of data, making them highly versatile. This means they can learn new tasks quickly with minimal examples, speeding up development.

## AI Connects with Tools

AI systems can now integrate smoothly with other software, databases, and online services. This allows AI to perform complex actions and automate real-world tasks.

## AI Understands Long Texts

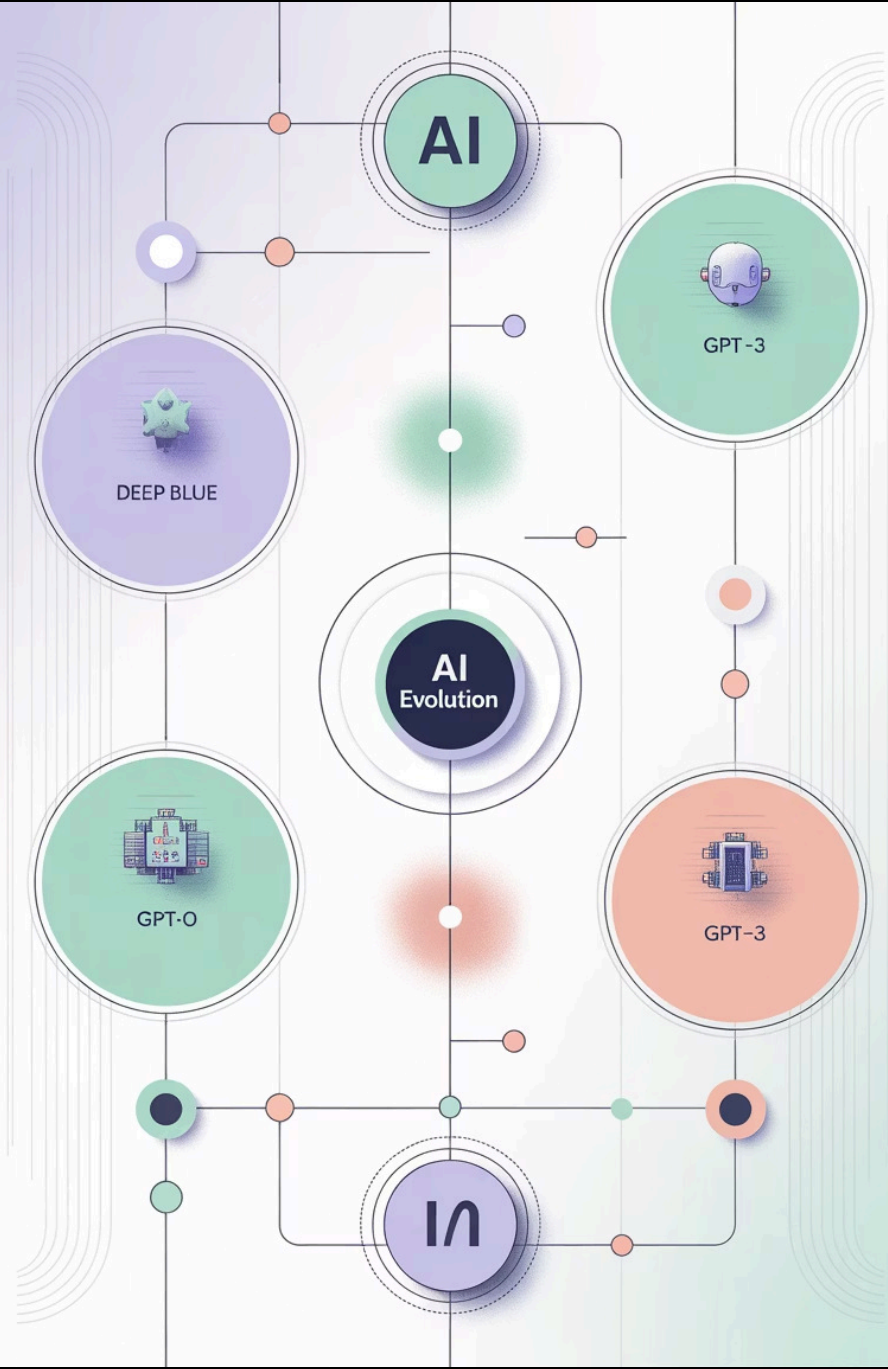
Modern AI can process and understand very long documents, not just short prompts. This allows for deep analysis of large texts, like research papers or books.

## Powerful AI for Everyone

Advanced AI models can now run efficiently on standard laptops. This makes powerful AI capabilities accessible to students and individuals without requiring specialized hardware.


## Rich Development Ecosystem

A comprehensive set of tools, open-source models, and widespread GPU access now accelerates AI development. This simplifies building and deploying AI applications.




# By the Numbers (2024–2025)

These key figures from 2024–2025 highlight AI's rapid growth and its transformative impact on industries, shaping new opportunities for professionals.



**Business AI Adoption**

A growing majority of organizations are actively using AI technologies in 2024, creating significant demand for skilled professionals. (AI Index 2025)



**Investment Growth**

Year-over-year increase in GenAI funding, reaching \$33.9B in private investment during 2024, fueling innovation and new roles.

The increasing accessibility of powerful AI, capable of running on standard laptops, is democratizing advanced capabilities. This shift, coupled with significant productivity gains in tasks like text generation and coding, means AI is reshaping how work gets done across diverse industries.

# Transformers vs. Traditional Architectures

While Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs) have been foundational in AI, the Transformer architecture has revolutionized how we process sequential data, especially in natural language.

## CNNs & RNNs: Specialized Processors

**CNNs:** Excel at capturing local patterns (e.g., image features). Process data in fixed-size windows.

**RNNs:** Handle sequential data by processing it step-by-step, maintaining a "memory" of previous steps.

- **Limitations:** RNNs struggle with long-range dependencies due to vanishing/exploding gradients. CNNs are less suited for long sequences without complex modifications.
- **Processing:** Predominantly sequential.
- **Key Uses:** CNNs for image classification, RNNs for speech recognition and simple text tasks.


## Transformers: Attention & Parallelism

**Self-Attention:** Allows the model to weigh the importance of different parts of the input sequence relative to each other, capturing global context.

**Parallel Processing:** Processes entire sequences simultaneously, significantly speeding up training and inference.

- **Strength:** Effectively handles very long-range dependencies and complex relationships within data.
- **Processing:** Highly parallelized.
- **Key Uses:** Natural Language Processing (NLP) tasks like translation, text generation, and increasingly, computer vision.

## Comparing architectures





# How AI Actually Works (LLMs & Transformers)

Ever wondered how AI models like ChatGPT can generate such human-like text and even code? It all boils down to two core concepts: Large Language Models (LLMs) and the Transformer architecture.

01

## Massive Data Ingestion

LLMs are trained on colossal datasets of text and code from the internet – billions of pages, books, articles, and conversations. They literally read the entire digital world.

02

## Pattern Recognition Engine

Their goal is to find statistical patterns in all that data. They learn which words tend to follow others, how sentences are structured, and even the subtle nuances of human conversation.

03

## The "Transformer" Architecture

This is the secret sauce! The Transformer is a special neural network design that allows the AI to process entire sequences of text simultaneously, understanding long-range dependencies and complex context within the data.

04

## Predicting the Next Word

When you give an LLM a prompt, its fundamental task is to predict the most probable next word in a sequence. It does this over and over, building coherent and contextually relevant responses, code, or stories.

05

## Sophisticated, Not Sentient

While incredibly powerful, LLMs don't "think" or "understand" in the human sense. They are highly advanced pattern-matching systems, generating text based on probabilities derived from their vast training data, not conscious thought.

# How AI Actually Works (LLMs & Transformers)

Modern AI, especially Large Language Models (LLMs) like ChatGPT, relies heavily on a brilliant design called the Transformer architecture. This innovative approach transformed how computers understand and generate text by focusing on how words relate to each other, rather than processing them one by one. Let's break down some of its core ideas with simple examples.

## 1. Self-Attention: Understanding Word Relationships

Imagine reading the sentence: "The **animal** didn't cross the street because **it** was too tired." As a human, you instantly know "it" refers to the "animal" and not the "street." Self-attention gives AI models a similar ability. It helps the model look at all the words in a sentence simultaneously and figure out how important each word is to understanding every other word.

Think of it like this: When the AI processes the word "it," it asks itself, "Which other words in this sentence are most relevant to 'it'?" It then assigns a 'score' to each word (animal, didn't, cross, street, because, was, too, tired) based on its relevance to "it." This allows the AI to correctly link "it" back to "animal," understanding the context much better.

## 2. Multi-Head Attention: Multiple Perspectives

If self-attention is like looking at a sentence to understand one type of relationship, Multi-Head Attention is like having several experts look at the same sentence, each from a different angle, all at once. Instead of just one 'relevance score' calculation, the model performs several self-attention calculations in parallel.

For instance, one "head" (expert) might focus on grammatical relationships (e.g., subject-verb agreement), another on semantic relationships (e.g., synonyms or related concepts), and another on sentiment. By combining insights from these multiple "heads," the AI gets a richer, more nuanced understanding of the input. It's like getting diverse opinions before making a decision, leading to a more accurate and comprehensive interpretation.

## 3. Positional Encoding: Keeping Words in Order

Unlike traditional methods that processed words sequentially, Transformers process all words in a sentence at the same time. While this is super-efficient, it means the model loses information about the order of words. For us, "Dog bites man" and "Man bites dog" mean very different things because of word order.

This is where Positional Encoding comes in. It's like adding a small, unique tag to each word that tells the model its exact position in the sentence. So, even though the words are processed in parallel, the model still "knows" which word came first, second, third, and so on. This ensures that the meaning derived from word order isn't lost, allowing the AI to differentiate between "Dog bites man" and "Man bites dog" just like we do.





## Positional Encoding: AI's Understanding of Order

While **Transformers** process words simultaneously for efficiency, this can lead to losing critical information about word order. Without it, "Dog bites man" could be misinterpreted as "Man bites dog."

Positional Encoding solves this by adding a unique "address tag" to each word, denoting its position in the sequence. This ensures the AI understands the grammatical structure and meaning derived from word order.

For example, "The cat sat on the mat" and "The mat sat on the cat" use the same words. Different positional encodings allow the AI to grasp their distinct meanings, mirroring human comprehension.

# Self-Attention: AI's Contextual Vision

Self-Attention is a revolutionary mechanism that allows AI models to weigh the importance of different words in an input sequence when processing each word. It's how the model builds a rich contextual understanding, similar to how we focus on key phrases in a sentence.

## → Connecting the Dots

Imagine the sentence: "The **animal** that I saw yesterday was huge." As humans, we instantly know "that" refers to "animal." Self-attention enables AI to make this same connection.

## → Every Word Looks at Every Other Word

For each word, the AI doesn't just process it in isolation. It simultaneously "looks" at all other words in the sentence and assigns an "attention score" to determine how relevant each word is to understanding the current word.

## → Building Contextual Meaning

By calculating these scores, the model understands that "that" is highly relevant to "animal" and much less so to "yesterday" or "huge." This creates a dynamic web of relationships, forming a nuanced contextual representation for each word.

This process is repeated for every word, creating a comprehensive understanding of the entire sentence's meaning, not just individual words. It's crucial for generating coherent and contextually appropriate text.

# Multi-Head Attention: Diverse Insights

Building on Self-Attention, Multi-Head Attention performs several 'attention calculations' in parallel. Imagine a team of experts meticulously analyzing the same sentence, each from a unique and specialized perspective. This parallel processing allows the AI to grasp a more nuanced and comprehensive understanding of the input.

<div>1</div> <div><b>Grammatical Expert</b><p>Focuses on sentence structure and word roles (e.g., subject-verb agreement), ensuring syntactic correctness and flow. For "The tired student submitted their assignment late...", this head would note "student" as the subject of "submitted" and "their" as the possessive pronoun relating to "student."</p></div>	<div>2</div> <div><b>Semantic Expert</b><p>Analyzes word meanings and conceptual relationships (e.g., synonyms, related concepts), extracting the core message. In "The tired student submitted their assignment late because it was challenging," this head would connect "tired" to the state of the "student," and "challenging" to the nature of the "assignment."</p></div>
<div>3</div> <div><b>Relational Expert</b><p>Identifies dependencies between words, like linking pronouns to their antecedents (e.g., "it" to "animal") for broader context. For the sentence "The tired student submitted their assignment late because it was challenging," this head would determine that "it" refers specifically to "assignment" and "their" to "student."</p></div>	<div>4</div> <div><b>Tonal Expert</b><p>Assesses the sentiment or emotional tone of the text (e.g., positive, negative, neutral), capturing subtle emotional cues. In "The tired student submitted their assignment late because it was challenging," this head would identify "tired," "late," and "challenging" as contributing to a sympathetic or slightly negative tone regarding the student's experience.</p></div>

By integrating these diverse viewpoints, the model synthesizes a richer, more robust representation of the text, crucial for generating contextually accurate and coherent language.



# How AI Actually Works (Simple Version)

Imagine AI as a very clever assistant that learns from vast amounts of information, much like we learn from reading books and listening to conversations. It's designed to recognize patterns and make predictions, simplifying complex tasks for us.

1

## 1. Breaking Down Information

Before AI can understand anything, it first breaks down all the text it processes into tiny, manageable pieces—like individual words or even parts of words. Think of it like organizing a library, where each word is a catalog entry.

2

## 2. Finding Connections

Once it has these tiny pieces, the AI doesn't just see them as separate. It figures out how they relate to each other, similar to how our brain connects 'apple' to 'fruit' and 'red'. This helps the AI understand the meaning and context of words.

3

## 3. Learning Patterns


This is where the AI "thinks" and finds patterns, much like a human learns grammar rules or the structure of a story. It's like a very smart autocomplete that has seen billions of examples and knows what usually comes next in a sequence.

4

## 4. Creating New Things

Based on all the patterns it has learned, the AI can then generate new text, summaries, ideas, or even images. It's like having a creative assistant that can "fill in the blanks" or write a whole new piece based on your prompt.

Understanding these basic steps helps students use AI more effectively by knowing that it learns from patterns and context, enabling them to craft better prompts and evaluate AI outputs.



# AI Engineering Ladder — Overview

Your AI engineering journey: A clear path from basic AI use to advanced system design, building crucial skills at each stage for practical application.

## Level 1: Using AI

Effectively communicate with AI (prompt engineering), and connect to AI tools for simple tasks and interactions.

## Level 2: Integrating AI

Enhance AI with relevant data, improve speed and efficiency, and safely integrate AI into existing systems and workflows.

## Level 3: Engineering Systems

Design robust AI applications with built-in safety, evaluate their performance, and combine multiple AI models for smarter solutions.

## Level 4: Optimizing at Scale

Deploy AI models efficiently, manage performance, monitor their behavior in real-time, and ensure data privacy and legal compliance.



# Choosing AI Models for Your Projects: A Student's Guide

## Proprietary Models (Cloud-Hosted)

These advanced, cloud-hosted models offer high performance, ease of use, and often provide free tiers or academic access. Ideal for quick prototyping and projects needing top-tier accuracy.

### Text & Code

- **GPT-4 class (OpenAI):** General intelligence & complex reasoning, advanced coding.
- **Claude 3.5 class (Anthropic):** Nuanced analysis, long context, creative writing.
- **Gemini 1.5+ (Google):** Multimodal understanding (text, image, video), powerful for diverse tasks.

### Vision

- **GPT-4o / Gemini Vision:** Integrated multimodal understanding for image analysis.
- **DALL-E 3 / Midjourney (Image Generation):** Creative asset generation, visual concept exploration.
- **Azure AI Vision / Google Cloud Vision AI:** Object detection, facial recognition, OCR (often with free tiers).

### Audio

- **OpenAI Whisper (Speech-to-Text):** Highly accurate transcription for various languages.
- **ElevenLabs / Google Cloud Text-to-Speech:** Realistic voice generation for narration, voiceovers.

When choosing a model for your student project, consider:

- **Project Needs:** Does it require cutting-edge performance (proprietary) or deep customization/privacy (open-source)?
- **Budget:** Leverage free tiers of proprietary services or fully open-source options to minimize cost.
- **Learning Goal:** Open-source models offer more hands-on learning with architecture and fine-tuning.
- **Resource Availability:** Proprietary models handle infrastructure; open-source might require local hardware or cloud setup knowledge.

Experiment with free access tiers for proprietary models and explore the vast open-source community for cost-effective and highly customizable solutions.

## Open-Source Models (Local/Cloud Deployment)

Community-driven models provide maximum flexibility, enhanced privacy, and cost control. Excellent for learning, custom fine-tuning, and projects requiring local deployment or specific hardware optimization.

### Text & Code


- **Llama 3.x family (Meta):** Scalable, powerful models for diverse tasks, fine-tunable.
- **Mistral / Mixtral:** Efficient and powerful, great for local or constrained environments.
- **Qwen 2/2.5 series (Alibaba):** General purpose and coding specialists, multilingual.
- **Phi-3 (Microsoft):** Compact models, suitable for edge computing and smaller projects.

### Vision


- **Stable Diffusion (Image Generation):** Highly customizable image generation, local control.
- **YOLO (You Only Look Once):** Real-time object detection and tracking.
- **LLaVA / Fuyu:** Multimodal (vision-language) understanding for image-text tasks.

### Audio

- **Wav2Vec / Whisper (Open-Source version):** Speech recognition for custom datasets.
- **Coqui TTS / VITS:** Text-to-speech for generating custom voices.







# Local vs Cloud Models — When to Use Which

## Local Models (On-Device)

### Privacy & Security

Complete data control with no external transmission, keeping sensitive information secure.

### Cost Efficiency & Speed

Zero per-token costs, low latency responses, and full offline functionality.

### Smaller Model Versatility

Smaller, optimized models can effectively handle many tasks locally, like chat and RAG.

Ideal for ideation, secure coding assistance, document summaries, and small-scale retrieval-augmented generation (RAG) applications.

## Cloud/Hosted Models

### Cutting-Edge Capabilities

Access to the latest, most powerful models for superior multimodal processing and complex tasks.

### Scalability & Management

Managed infrastructure handles scaling, maintenance, and updates, reducing operational overhead.

Best suited for high-demand applications, advanced AI research, and situations requiring extensive computational resources or diverse model types.

**Key Takeaway:** Choose local for privacy, cost control, and basic tasks; opt for cloud when state-of-the-art capabilities, scalability, and managed infrastructure are critical.



# AI Coding IDEs & Tools (2025)

## Free Local Setup

Get started with VS Code + Continue extension using Ollama. Recommended models include phi3:mini, llama3.2:3b-instruct, and mistral:7b-instruct for local, private AI coding.



## Hosted Platforms

Utilize platforms like Claude Projects for in-depth conversations, Windsurf for collaborative coding, and specialized tools (Lovable, Bolt.new, v0.dev) for rapid prototyping and generating boilerplate code quickly.



## IDE Integration

Boost productivity with AI-native IDEs like Cursor, seamless GitHub Copilot integration, Codeium (a free alternative), enterprise-grade Amazon Q, or JetBrains AI for their ecosystem. These tools provide AI assistance directly in your workflow.



## Actionable Best Practices

Paste failing tests for contextual debugging, request small, focused code changes, and keep prompts concise and specific to achieve optimal results from your AI coding assistant.

# Making AI Smarter with Your Own Information

Imagine you could teach an AI everything you know, specifically from your own documents, class notes, or company reports. That's exactly what Retrieval-Augmented Generation (RAG) does! It's like giving a powerful AI access to your personal knowledge base, allowing it to answer questions accurately using only the information you trust, instead of just what it learned from the internet.

1

## 1. Prepare Your Documents

First, we take your documents (like your lecture notes or company handbooks) and break them down into small, meaningful pieces. This helps the AI easily find the exact information it needs later.

2

## 2. Make Them Searchable

Next, these pieces of information are converted into a format that the AI can quickly search through. Think of it like creating a super-fast, intelligent index for all your content, so the AI can quickly match what you're asking to what you've provided.

3

## 3. Find the Best Match

When you ask a question (e.g., "What was covered in last week's history class?"), the RAG system rapidly sifts through your indexed information to find the most relevant document pieces. It then double-checks to make sure it picks the very best and most accurate information.

4

## 4. AI Answers from Your Data

Finally, the AI uses these specifically found pieces of your information as a direct reference to formulate a precise and well-supported answer. This way, the AI isn't guessing; it's giving you an answer directly grounded in your specific data.

This method ensures the AI provides reliable, up-to-date answers based on your unique information, making it incredibly useful for personalized learning, business operations, and ensuring factual accuracy in specific contexts.



# Keeping AI Safe and Reliable

Using AI effectively means understanding how to get the most out of it while staying safe. Here are some simple tips to help you use AI tools responsibly and confidently every day.



## Always Double-Check Facts

AI can sound very confident, even when it's wrong. For important decisions or information, always cross-reference AI's answers with reliable sources, just like you would with any new information.



## Guard Your Personal Info

Be careful about what you share. Avoid giving AI sensitive personal details like passwords, bank account numbers, or private medical information. Assume anything you tell an AI could potentially be learned and used later.



## Watch for Biased or Strange Answers

AI learns from data, which can sometimes reflect biases or give unexpected responses. If an answer seems unfair, strange, or too good to be true, question it. The AI might have been "tricked" or trained on skewed information.



## Give Clear Instructions

To get the best from AI, be specific and clear with your questions. If an answer isn't quite right, try rephrasing your prompt or providing more context. Think of it like training a new assistant – the clearer your directions, the better the outcome.

# Prompt Patterns (Level 1 — Using AI)

Master these fundamental prompt patterns to dramatically improve AI output quality and reliability across all your applications.



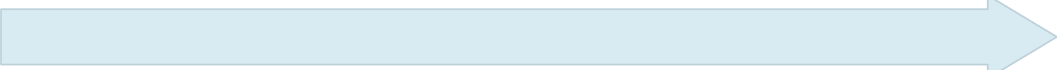
## Zero-Shot Prompting

Give clear instructions directly. For simple tasks, like "Summarize this article in 3 bullet points," it's effective without examples.



## Few-Shot Learning

Provide 2–3 high-quality examples showing the desired input-output format. This significantly boosts consistency and quality for complex tasks.




## Chain-of-Thought

Request step-by-step reasoning before the final answer. Using phrases like "Let's think through this step by step" improves AI's problem-solving.




## Self-Critique

Ask the AI to check for missing assumptions, potential risks, or alternative perspectives. This enhances output reliability and completeness.

 **Pro Tip:** Combine patterns for maximum effectiveness. For complex analytical tasks, integrate few-shot examples, chain-of-thought reasoning, and self-critique.


# Prompt Examples You Can Use Right Now

This slide shows practical, copy-paste ready prompts for common student tasks.




**Study Helper prompts**

- Explain [topic] like I'm 12 years old, then give me 3 real-world examples and 2 practice questions.
- Create a study guide for [topic] with key concepts, definitions, and memory tricks.
- Turn these messy notes into a clear summary: [paste notes].




**Writing Assistant prompts**

- Improve this paragraph for clarity and engagement: [paste text].
- Help me write a compelling introduction for an essay about [topic].
- Check this email for tone and professionalism before I send it: [paste email].




**Research prompts**

- Give me 5 different perspectives on [controversial topic] with credible sources.
- Create a research outline for [topic] with 10 key questions I should investigate.
- Fact-check this claim and provide counter-evidence: [paste claim].




**Creative prompts**

- Generate 10 unique project ideas for [subject] that would impress professors.
- Help me brainstorm creative solutions to this problem: [describe problem].
- Create 5 different story concepts based on this theme: [theme].




**Problem-solving prompts**

- Break down this complex problem step-by-step: [problem].
- What are 3 different approaches to solve [challenge] and their pros/cons?
- Help me debug this issue by asking clarifying questions: [describe issue].



**Power User Tips**

- Combine multiple prompt patterns (e.g., Chain-of-Thought with Few-Shot).
- Use follow-up questions to refine and clarify AI responses.
- Iterate on responses: ask for specific changes or alternative formats.
- Define clear constraints like desired length, tone, or output format.
- Request sources or justifications for factual claims made by the AI.

 Students should always fact-check AI responses and use them as starting points, not final answers.



# Ethical AI: Core Principles


Integrating ethical considerations throughout the AI lifecycle ensures responsible and trustworthy systems.

## Honest Uncertainty

Models should admit when they "don't know" rather than fabricating answers, preventing misinformation and building trust.

## Source Attribution

Always attribute sources for factual claims to ensure transparency and allow users to verify information.



## Privacy Protection

Safeguard user data by minimizing sensitive inputs, redacting information, and adhering to privacy regulations.

## Injection Defense

Protect against malicious prompt injections by validating inputs and restricting unauthorized commands.

# Interactive: AI Career Map (15m)

## Your Personalized Career Journey

This cool exercise helps you figure out what AI jobs are out there, what skills you'll need, and how to get from where you are now to where you want to be in the exciting world of AI!

01

### What Actually Interests You?

Let's figure out what really sparks your interest in AI! Do you love **building cool stuff**, diving deep into **data analysis**, or using AI to **solve real-world problems** and help people? Think about what you enjoy doing, and what kind of work environment you thrive in.

02

### Discover Your AI Superpowers (and Jobs!)

Now that you know what you like, let's explore some actual entry-level AI roles! Think **Junior AI Developer**, **Data Analyst Intern**, **AI Research Assistant**, or **Machine Learning Engineer Trainee**. What do these jobs actually involve? What's a typical day like? Where can they lead you?

03

### What Skills Do You Need?

You've got some skills already, awesome! Now, compare what you know with what these cool AI jobs actually need. This helps you spot any **skill gaps** and figure out what to learn next to crush your career goals.

04

### Build Your Learning Game Plan

Ready to level up? Let's create a practical **2-week sprint** action plan! This includes fun projects, helpful resources, and mini-milestones to help you grab those AI career goals. We're talking bite-sized learning!

05

### Save Your AI Adventure Map

Don't let your hard work disappear! Document everything you've learned and discovered. This will be your personal treasure map for ongoing career development, ready whenever you need it.

# College Projects — Campus Life

Transform your daily college experience with practical AI applications that solve real problems and save time.



## Attendance Tracker

Digitize handwritten attendance to CSV and Google Sheets for simple, automated tracking.



## Timetable Optimizer

Optimize your schedule: avoid conflicts, maximize free time, and find ideal course combinations.



## Event Marketing

Generate eye-catching posters and compelling promotional copy for campus activities.



## Study Assistant

Convert lecture notes into summaries, custom quizzes, and flashcards for effective studying.



## Campus Issues Bot

Auto-draft professional emails for complaints, track resolutions, and streamline admin communication.



## Lecture Insights

Analyze recorded lectures to generate timestamped summaries, key concepts, and study guides.

# College Projects — Placement & Coding

Accelerate your career preparation with AI tools designed specifically for technical interviews and skill development.

1

## Resume Bullet Generator

Generate impactful, metrics-driven resume bullets. Highlight your achievements to grab recruiter attention instantly.

2

## Mock Interview Simulator

Practice with AI-generated interview questions. Get instant feedback and build confidence through realistic conversations.

3

## DSA Learning Coach

Master Data Structures & Algorithms with personalized practice. Get strategic hints, test cases, and track your progress to ace coding challenges.

4

## Code Review Assistant

Enhance code quality with AI-powered review. Get automated analysis, suggestions for test cases, and learn best practices.

5

## System Design Mentor

Design robust systems with AI. Generate architectural diagrams, identify potential bottlenecks, and improve scalability with expert critiques.



# College Projects — Creators & Communities

Amplify your creative output and community engagement with AI-powered content creation and management tools.



## YouTube Optimization

Boost video discoverability and engagement with AI-generated chapters and SEO tags.



## Stream Highlights

Capture and share live stream highlights instantly for social media.



## Meme Generator

Design viral memes and captions using trending templates.

## Community Assistant Bot

Automate community support with bots for FAQs, forms, and teaching assistance.

## Personal Portfolio Builder

Build professional personal websites to showcase your projects and skills.

# What to Measure (for demos)

To demonstrate the real impact of your AI projects, focus on these practical metrics during your demos.



## Time Saved

Track time saved per task and measure active user engagement.



## Output Quality

Assess correctness, relevance, and overall success rate of generated content.



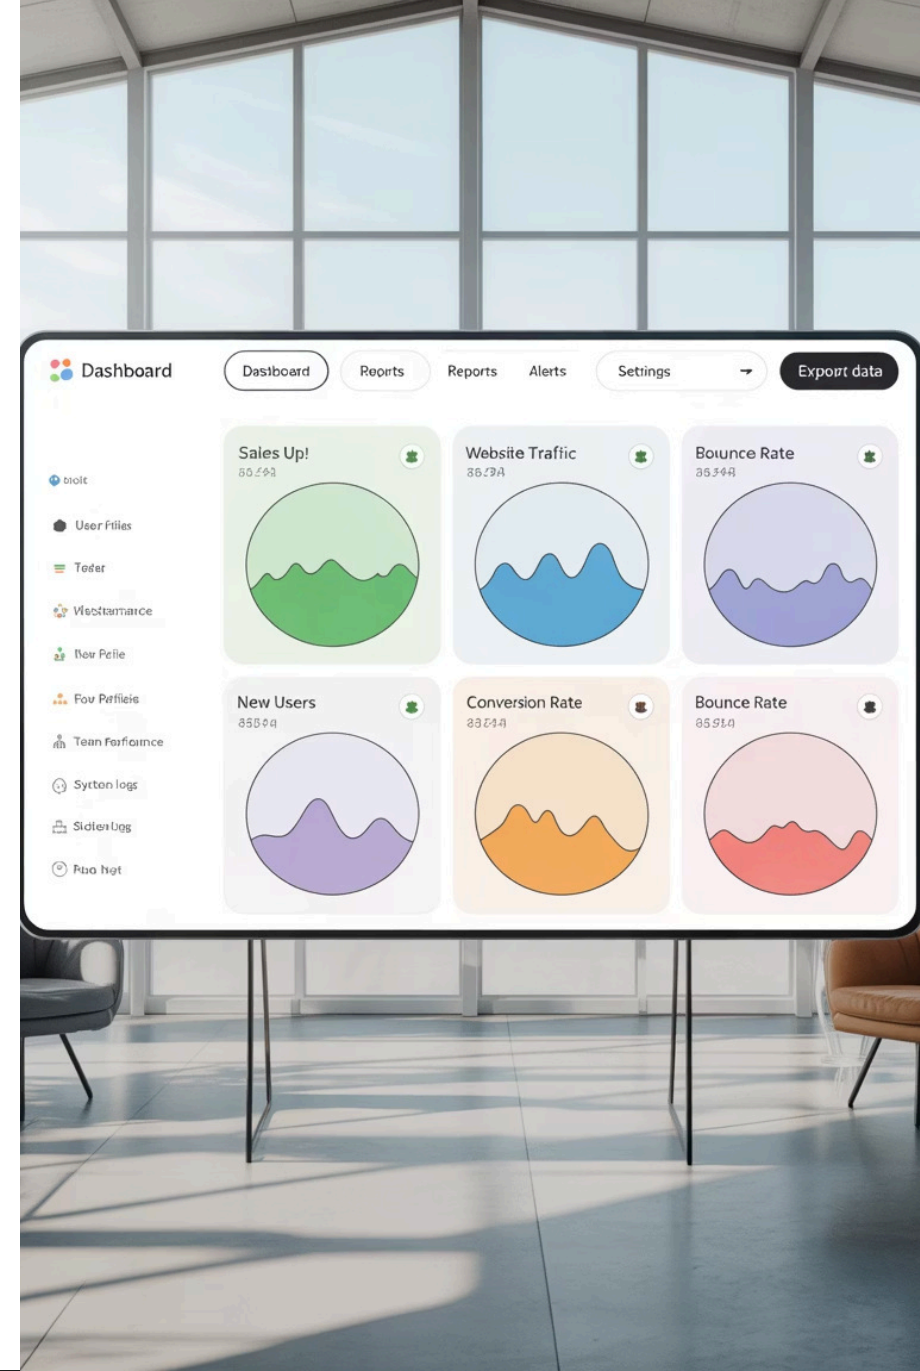
## Production Speed

Quantify content creation: posts, videos, or summaries generated per hour.



## Efficiency & Cost

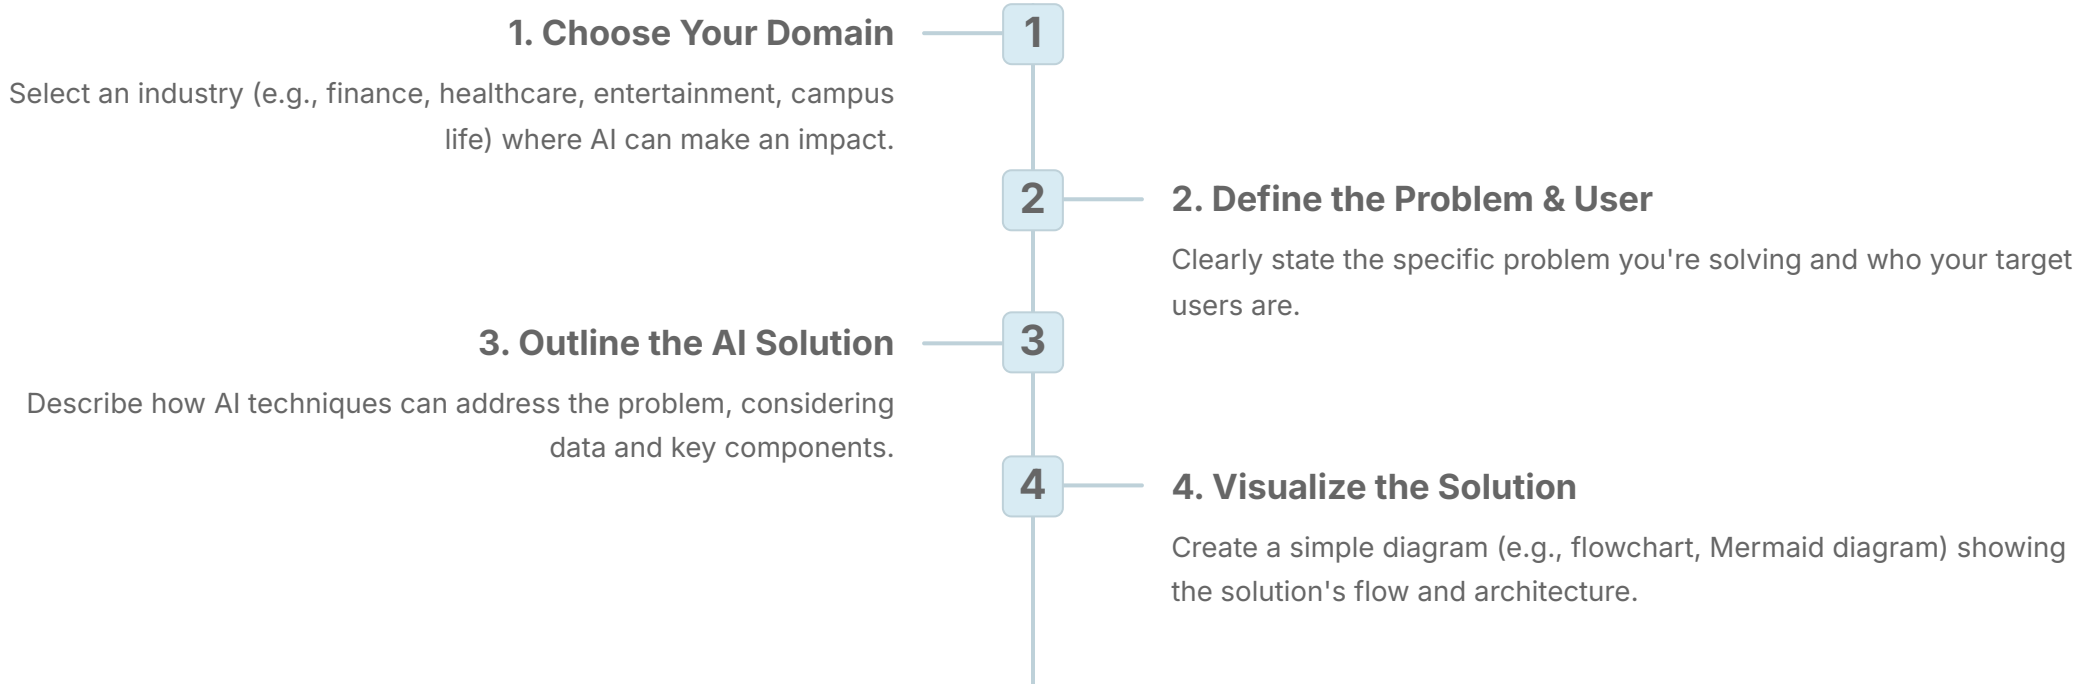
Monitor operational costs, responsiveness, and user satisfaction.





# Activity: Design an AI Use Case

Apply what you've learned to design a practical AI solution. Work in pairs to create a comprehensive use case for a real-world problem.



## Deliverable

Prepare one slide or note summarizing your project: **Problem** → **Approach** → **Success Metric** + **Technical Diagram**. Be ready to share your innovative solution!